




# Introduction Prompt Engineering I

Enabling Effective Interaction with AI  
systems



# A Definition of Prompt Engineering



- ❖ **Definition:** Prompt engineering involves crafting specific inputs (prompts) to guide AI models in generating desired outputs.
  - ❖ **Purpose:** It enhances the relevance and accuracy of AI responses by providing clear and structured prompts.
  - ❖ **Application:** Utilized in AI tasks like text summarization, translation, and content generation to achieve optimal results.
- 
- ❖ **Discussion:** What techniques do you know about or use daily with LLM interactions?



# Terminology





# Tokens



## ❖ Definition:

- Tokens are the basic units of text that language models process.
- Often fragments of words, punctuation, or individual letters.
- By breaking text into tokens, a model can handle variable input lengths efficiently. Query: “The candidate was taken aback by the interviewer’s question.”
- Tokenized Query: "The", "cand", "idate", "was", "taken", "aback", "by", "the", "interviewer", "", "s", "question", ".“

## ❖ Questions:

- Do LLMs like OpenAI ChatGPT and Google Gemini use tokens to remove sensitive information like phone numbers or salaries that may be accidentally prompted?
- How does tokenization vary across languages?






# Embeddings

## ❖ Definition:

- Embeddings are numerical representations of words, phrases, or sentences that capture their meaning in a multi-dimensional space.
- These vectorized forms help models understand semantic relationships between words, even if they are not explicitly stated.
- Embeddings enable tasks like similarity detection, topic modeling, and retrieval-based AI.

## ❖ Questions:

- How do embeddings help LLMs understand words that weren't in their training data?
- Why are embeddings useful for tasks like search, recommendation systems, or text classification?






# Training



## ❖ Definition:

- Training in a large language model is the process of teaching the model to understand and generate text. In simple terms, this involves feeding the model a large dataset of text so it can learn patterns, grammar, and context.
- Key hyperparameters are used to fine-tune this learning process:
  - Learning rate (how fast the model updates its weights)
  - Number of layers (the network's depth),
  - Batch size (how many examples are processed at once)
  - Temperature, which affects response creativity, is used during *inference*.



## ❖ Questions:

- Why do you think increasing the amount of training data positively affects a model's performance?
- What do you think is a key challenge with respect to training?



# Overfitting



## ❖ Definition:

- Overfitting occurs when a model learns noise and random fluctuations in training data rather than the underlying pattern. This leads to excellent performance on training data but poor generalization to new data.

## ❖ Key Points:

- Caused by excessive model complexity relative to the amount of training data.
- Results in high variance, where training accuracy is high, but test accuracy is low.
- Mitigation methods include regularization, early stopping, and cross-validation.

## ❖ Questions:

- How can you detect overfitting in a model's performance?
- Why do LLMs require a large amount of data?





# Inference



## ❖ Definition:

- Inference in a large language model is the process of generating text from a pre-trained model. In simple terms, when you enter a prompt (like a question), the model responds with an answer.
- During inference, text is converted into tokens.
- The LLM processes these tokens to predict the next token based on the previous ones. In other words, predict the next word.
- This word is then added to the sequence and the process repeated.

## ❖ Questions:

- What role do tokens play during inference?
- Why do LLMs sometimes produce incorrect or nonsensical outputs?





# Context Window



## ❖ Definition:

- The context window is the **maximum** number of tokens a language model can process at once, including both input (prompt) and output (response).
- If the total tokens exceed this limit, older tokens are truncated, which can lead to incomplete responses or loss of important information.
- Managing the context window effectively is crucial for tasks requiring long conversations, document analysis, or multi-step reasoning.

## ❖ Questions:

- If I ask an LLM "How are you" and it responds, how many tokens are used from the context window?
- If I then ask a follow-up question "What is  $1 + 1$ " how many tokens are now used from the context window?
- What happens when the window is exceeded?






# Temperature

## ❖ Definition:

- Temperature controls the randomness of an LLM's responses by adjusting the probability distribution of word selection.
- A higher temperature value makes the output more diverse and creative by allowing the model to pick less likely words, while a lower temperature makes responses more focused and deterministic by favoring the most probable words.

## ❖ Questions:

- How should you set the temperature for a science-based question?
- Can I set the temperature in a normal prompt window?






# Fundamental Principles





# Fundamentals (Technique-agnostic)





# Structure Prompts



## ❖ Include distinctions: Clearly differentiate between


- Instruct
- Context
- Tone
- Detail
- Format

## ❖ Straightforward example:

- Translate "Hello, world!" into French.

## ❖ More complex examples require more explanation:

- Compare Tesla Model Y and BMW iX2. Present your answer as a table using metric units and including key attributes like price (in SGD), range (km), and acceleration (0–100 km/h).





# Specify Output Format




## ❖ Clearly specify the expected output format:

- Paragraph
- Code
- Table
- Brief or Detailed
- Bullet points or enumerated lists
- JSON, HTML or Other Format

## ❖ Clearly stating the format helps the model generate the desired results.

## ❖ Example, "Give me a list of the top 5 tourist attractions in Paris."





# Evaluate



- ❖ **Define Evaluation Criteria:** Consider clarity, correctness, completeness, and format adherence. → What aspects of the response are most important to evaluate?
- ❖ **Confirm Intent Alignment:** Check if the response followed instructions, answered the question, and stayed relevant. → Did the response understand and address the core request?
- ❖ **Assess Logic and Accuracy:** Examine the reasoning and factual correctness. → Is the response logically sound and factually accurate?





# Evaluate



- ❖ **Check Technical Details:** Verify consistency and correctness of numeric and technical data. → Are the technical details presented accurate and consistent within the response?
- ❖ **Determine Scrutiny Level:** Decide on the depth of evaluation needed based on query complexity. → Does this query require a quick check or a more in-depth evaluation?





# Experiment




- ❖ Experimentation: An **observational and iterative** approach to Prompt Engineering is absolutely necessary.
- ❖ Evidence-based Process: Success relies on real-world **testing and iteration** and evaluation.
- ❖ Diverse Approaches: Experiment with **variations** in wording, format, and context.
- ❖ Continuous Learning: **Refine** your prompts based on feedback and observed results.
- ❖ Adaptability: **Tailor** your prompts to different tasks and evolving requirements.






# Fundamentals (Technique-agnostic)





# Prompt Engineering Techniques 1





# Prompt Engineering Techniques (Part 1)

## ❖ Zero-shot

- Your prompt contains no context and no examples.
- The model infers intent from the prompt alone.
- Works well with straightforward queries, where minimal guidance is required.

## ❖ Examples:

- Explain how photosynthesis works.
- Translate "Goodbye cruel world" into French.
- Who was the President of the USA on 21 January 2025?

?

Will such queries always work reliably?

## ❖ Side-effects:

- Vague or incorrect responses
- Relies almost completely on the model's internal knowledge
- Consistency can vary significantly.
- The model has no guidance for internal reasoning.

# Prompt Engineering Techniques (Part 1)



## ❖ One-shot

- Providing an example can help shape output content and format.
- Useful when a **single** illustration can clarify style and/or format of output.
- Enables models to understand and perform accurately when extensive context is not required.
- One-shot prompting is a form of **few shot** prompting, but with a **single** example.



## ❖ Examples:

- **Determine the sentiment** in this review. Example: "Exceeded my expectations" -> Positive
- **Correct the grammar** in this article. Example: "She don't like apples." → "She doesn't like apples."
- **Classify this article** as news, opinion or ad. Example: "The stock market reached an all-time high today." → News

## ❖ Side-effects:

- Model may struggle with complex tasks requiring deeper understanding or extensive context.
- A poorly written prompt or example may lead to inaccurate or suboptimal output.
- Overfitting is a possibility.



# Prompt Engineering Techniques (Part 1)

## ❖ Few-shot

- Offers multiple examples to guide the response
- Improves accuracy and consistency for more complex tasks
- The model learns style, tone, and format from these examples



## ❖ Examples:

- Here are three examples of product descriptions:
  - Example: Lenovo ...
  - Acer ...
  - Asus...
  - Write a similar product description for a MacBook Pro M4.
- Classify the following as 'spam' or 'not spam'.
  - Example: "You've won a free vacation. Click here to claim your prize." Classification: **Spam**
  - Text: "Meeting tomorrow at 10 AM in the conference room." Classification: **Not Spam**
  - Text: "Limited-time offer! Get 50% off on all products." Classification: **Spam**
  - To Classify: "Your package has been shipped and will arrive by Monday."



# Prompt Engineering Techniques (Part 1)

## ❖ Few-shot (cont'd) - Side-effects

- **Example Sensitivity:** The model's performance heavily depends on the quality and selection of examples provided
- **Complex Reasoning:** As task/example complexity increases, few-shot prompting may still struggle to with deep reasoning as models may not grasp intricate relationships expressed in the examples
- **Overfitting Risk:** Providing too many examples can lead to overfitting



# Prompt Engineering Techniques (Part 1)




## ❖ Self-Critique (or Self-Evaluation)

- **Self-Critique** is a prompting technique where an AI model evaluates its own responses to identify errors, inconsistencies, or areas for improvement, and then refines its output accordingly.

## ❖ Steps

- **Initial Response Generation:** The AI provides an initial answer to a user's prompt
- **Self-Critique:** The AI reviews its own response, assessing accuracy, relevance, coherence and completeness
- **Issue Identification:** The AI identifies any errors, gaps, omissions or areas that require enhancement
- **Revision:** Based on its own evaluation, the AI revises its original response
- **Output:** The improved, refined response is presented to the user



## ❖ Reasoning Models

- It is important to understand that reasoning models, like OpenAI o1 and DeepSeek R1 employ chain-of-thought (CoT) reasoning and self-reflection. These techniques are a form of self-critique.



# Prompt Engineering Techniques (Part 1)

## ❖ Self-Critique (or Self-Evaluation) cont'd

## ❖ Examples

- Explain the process of photosynthesis. Then, review your explanation for accuracy and completeness, and provide any necessary improvements.
- Detail the process of nuclear fusion. Afterwards, evaluate your response for its thoroughness and precision, and suggest improvements where needed.
- Write a function in Python that sorts a list of integers. Then, critique your code for efficiency and readability, and suggest enhancements.

## ❖ Side-effects


- Using self-critique prompts may produce
  - shallow or generic self-critiques,
  - self contradictions
  - repetitive or off-target “improvements”



# Summary



**Zero-shot**  
No prior examples. Model infers from the prompt intent.  
Ok for straightforward queries.



**Few-shot**  
Multiple examples provided.  
Improvement in accuracy and consistency for complex tasks. Overfitting risk.

**One-shot**  
Provide one example. The example can shape output and format. Ok when extensive context is not required.

**Self-critique**

Encourages iterative improvement through self-assessment.



# Evaluating Prompt Output

**Activity**



# Activity



- ❖ Open the document called Prompt Creation Activity
- ❖ Your team number determines your area of investigation:
  - Sentiment analysis
  - Summarisation
  - Translation
  - Classification
  - Change of Tone
- ❖ Follow the timing given in the document
- ❖ Present findings



# Discussion Questions

- ❖ How did the zero-shot, one-shot, and few-shot prompts compare in terms of clarity and detail?
- ❖ Which type of prompt produced the most accurate or creative response?
- ❖ Did adding self-critique change the quality of the output? In what way?
- ❖ What surprised you the most about the model's responses?
- ❖ How might you refine your prompts in the future for better results?
- ❖ After ranking each prompt and output from 0–5 and calculating the correlation across all rubric categories, what did you learn about how prompt quality affects output quality?

AI



# Keeping track

**Activity**





# Activity

- ❖ Open the document called Note Taking Activity.
- ❖ Follow the timing
- ❖ Complete the activity.

Ai



# Lesson 02

**Summary of Material Covered**

# Lesson 02 Review



- ❖ Overview of Prompt Engineering 1:
  - **Tokens** – Basic units of text (words, sub-words, punctuation) used by AI.
  - **Embeddings** – Numeric representations capturing word relationships.
  - **Training** – Teaching AI models using large datasets.
  - **Overfitting** – Model learns noise instead of patterns, reducing generalization.
  - **Inference** – AI generates text based on learned knowledge.
  - **Context Window** – Limit on the number of tokens AI can process at once.
  - **Temperature** – Controls randomness in AI responses (low = precise, high = creative).
- ❖ Structuring Prompts
  - **Instruction:** Tell the AI what you want it to do.
  - **Context:** Give the AI background information to better understand your request.
  - **Tone:** Specify the style or attitude you want the AI to adopt in its response.
  - **Detail:** Include specific points or elements you want the AI to address.
  - **Format:** Define the structure or layout for the AI's response.



# Lesson 02 Review

## ❖ Prompting techniques:

- **Zero-shot** – No examples, AI infers intent
- **One-shot** – Single example for guidance (version of few-shot)
- **Few-shot** – Multiple examples improve accuracy
- **Self-Critique** – AI evaluates and refines its own output

Ai



# To Do (Homework)

- ❖ **Clear a minimum of 10Gb** of hard disk space on your computer. If you need to visit the IT helpdesk to help you clear, do so asap.
- ❖ Follow the instructions in the **LM Studio Install** document which is in the **EdTech Resources** folder.
  
- ❖ **Please Note: Downloading large language model files over the RP network is very slow. Please do this task at home.**

Ai

# Thank you

School of Infocomm

C240 AI Essentials and Innovations

© Republic Polytechnic 2025: All Rights Reserved